

Heart Disease Dataset Project

April 14, 2022

```
[74]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[92]: df = pd.read_csv("heart.csv")
```

```
[93]: df.head()
```

```
[93]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[94]: df.shape
```

```
[94]: (303, 14)
```

```
[96]: df.describe()
```

```
[96]:
```

	age	sex	cp	trestbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	

50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000

	restecg	thalach	exang	oldpeak	slope	ca \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[97]: df.isnull().sum()
```

```
[97]: age          0
sex            0
cp             0
trestbps      0
chol          0
fbs           0
restecg       0
thalach       0
exang         0
oldpeak       0
slope         0
ca            0
thal          0
target        0
dtype: int64
```

```
[98]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	303 non-null	int64
1	sex	303 non-null	int64
2	cp	303 non-null	int64
3	trestbps	303 non-null	int64
4	chol	303 non-null	int64
5	fbs	303 non-null	int64
6	restecg	303 non-null	int64
7	thalach	303 non-null	int64
8	exang	303 non-null	int64
9	oldpeak	303 non-null	float64
10	slope	303 non-null	int64
11	ca	303 non-null	int64
12	thal	303 non-null	int64
13	target	303 non-null	int64


dtypes: float64(1), int64(13)

memory usage: 33.3 KB

None

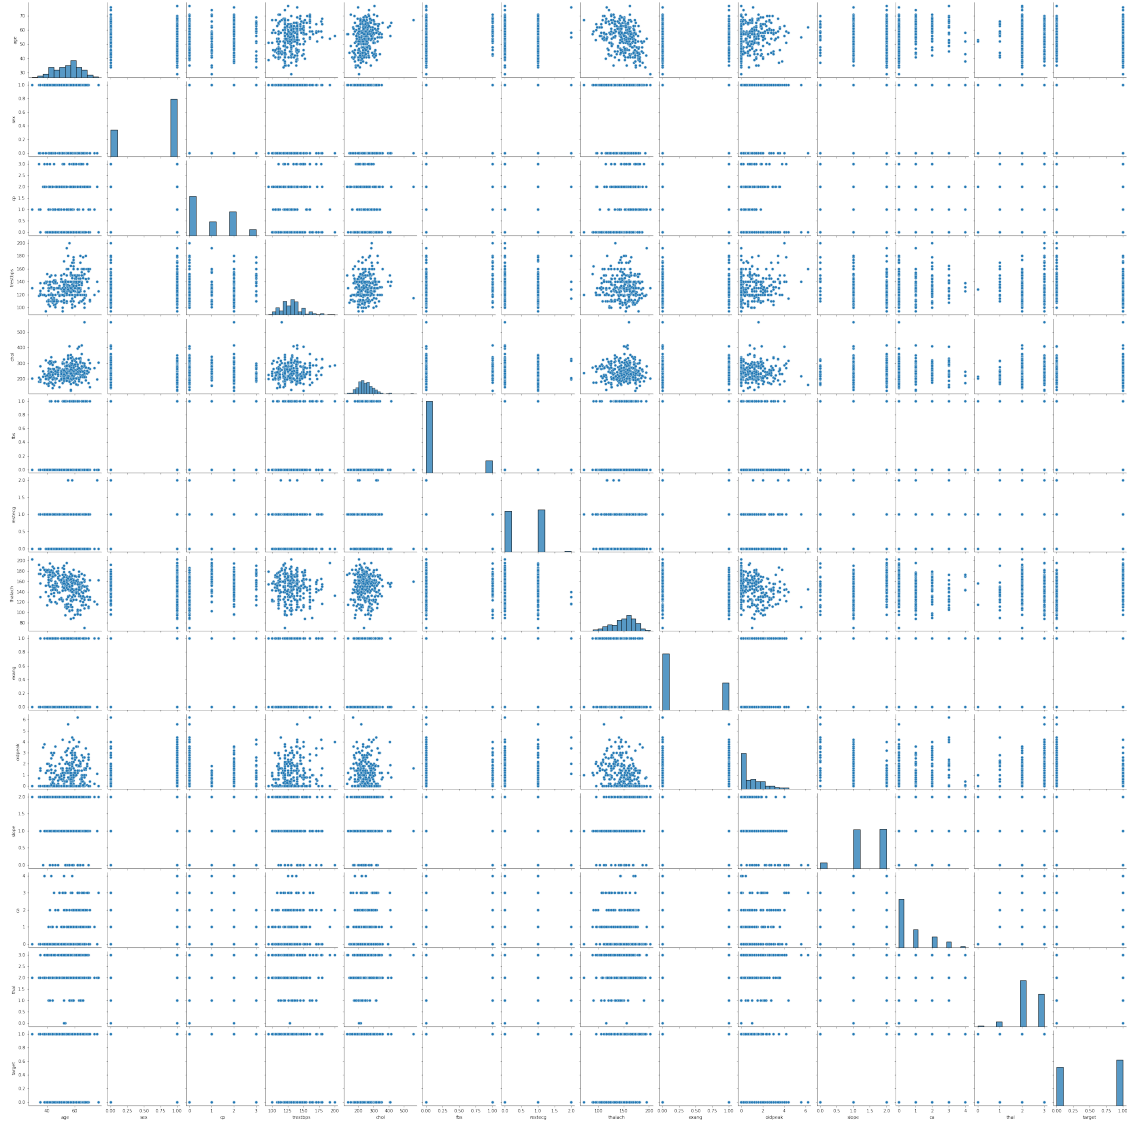
```
[99]: plt.figure(figsize=(20,10))
      sns.heatmap(df.corr(), annot=True, cmap='terrain')
```

[99]: <AxesSubplot:>




```
[101]: sns.pairplot(data=df)
```

```
[101]: <seaborn.axisgrid.PairGrid at 0x1f9501cb7f0>
```




```
[102]: sns.countplot(x='ca',hue='target',data=df)
```

```
[102]: <AxesSubplot:xlabel='ca', ylabel='count'>
```



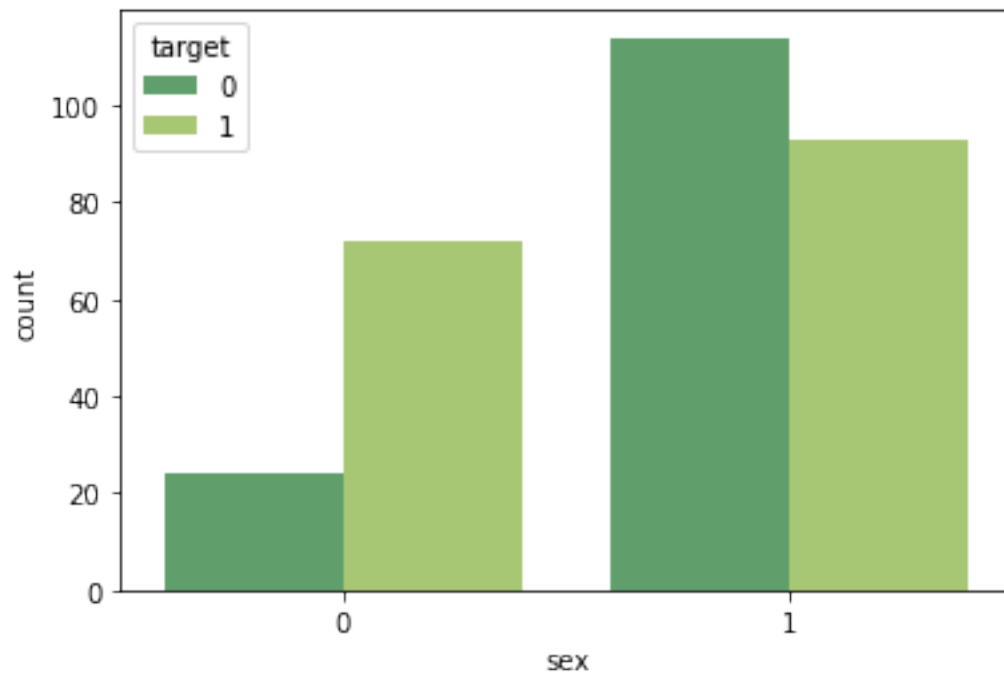
```
[105]: sns.barplot(data=df, x='sex', y='chol', hue='target', palette='summer')
```

```
[105]: <AxesSubplot:xlabel='sex', ylabel='chol'>
```




```
[113]: sns.countplot(x='sex', data=df, palette='summer', hue='target')
```

```
[113]: <AxesSubplot:xlabel='sex', ylabel='count'>
```




```
[119]: sns.countplot(x='target', palette='BuPu', data=df)
```

```
[119]: <AxesSubplot:xlabel='target', ylabel='count'>
```



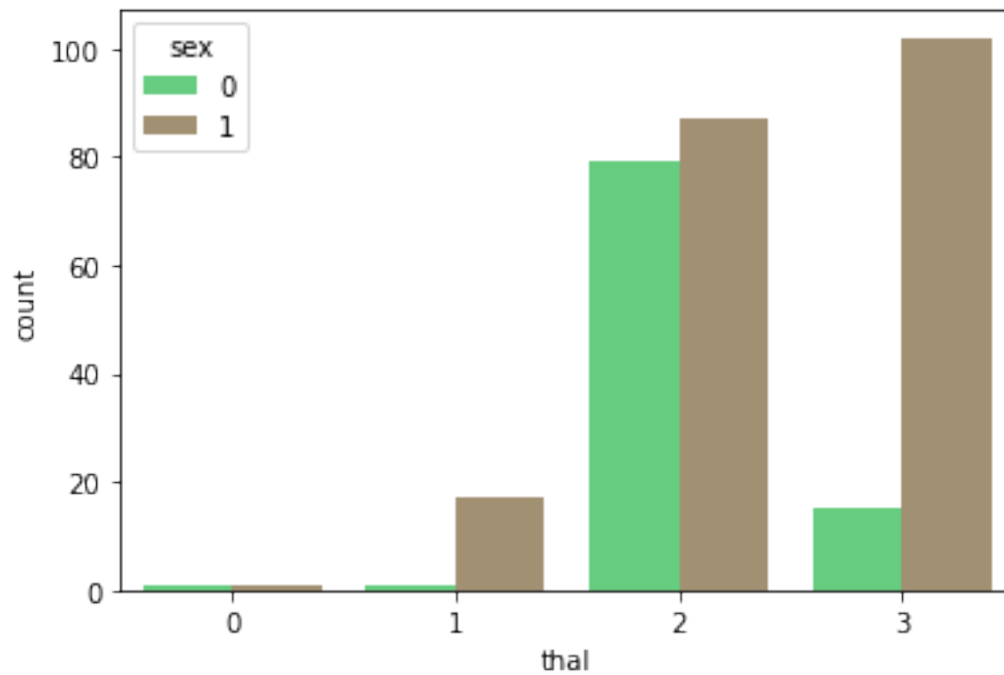
```
[108]: sns.countplot(x='thal',data=df, hue='target', palette='BuPu' )
```

```
[108]: <AxesSubplot:xlabel='thal', ylabel='count'>
```




```
[109]: sns.countplot(x='thal', hue='sex', data=df, palette='terrain')
```

```
[109]: <AxesSubplot:xlabel='thal', ylabel='count'>
```




```
[110]: sns.countplot(x='cp', hue='target', data=df, palette='rocket')
```

```
[110]: <AxesSubplot:xlabel='cp', ylabel='count'>
```

```
[111]: sns.countplot(x='cp', hue='sex', data=df, palette='BrBG')
```

```
[111]: <AxesSubplot:xlabel='cp', ylabel='count'>
```



```
[120]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
StandardScaler = StandardScaler()
columns_to_scale = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
df[columns_to_scale] = StandardScaler.fit_transform(df[columns_to_scale])
```

```
[121]: df.head()
```

```
[121]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	\
0	0.952197	1	3	0.763956	-0.256334	1	0	0.015443	0	
1	-1.915313	1	2	-0.092738	0.072199	0	1	1.633471	0	
2	-1.474158	0	1	-0.092738	-0.816773	0	0	0.977514	0	
3	0.180175	1	1	-0.663867	-0.198357	0	1	1.239897	0	
4	0.290464	0	0	-0.663867	2.082050	0	1	0.583939	1	

	oldpeak	slope	ca	thal	target
0	1.087338	0	0	1	1
1	2.122573	0	0	2	1
2	0.310912	2	0	2	1
3	-0.206705	2	0	2	1
4	-0.379244	2	0	2	1

```
[122]: X= df.drop(['target'], axis=1)
y= df['target']
```

```
[123]: X_train, X_test,y_train, y_test=train_test_split(X,y,test_size=0.
↪3,random_state=40)
```

```
[124]: print('X_train-', X_train.size)
print('X_test-',X_test.size)
print('y_train-', y_train.size)
print('y_test-', y_test.size)
```

```
X_train- 2756
X_test- 1183
y_train- 212
y_test- 91
```

Model Logistic Regression

```
[126]: from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()

model1=lr.fit(X_train,y_train)
prediction1=model1.predict(X_test)
```


```
[127]: from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,prediction1)
cm
```

```
[127]: array([[36,  4],
        [ 3, 48]], dtype=int64)
```

```
[129]: sns.heatmap(cm, annot=True,cmap='BrBG')
```

```
[129]: <AxesSubplot:>
```



```
[130]: TP=cm[0][0]
TN=cm[1][1]
FN=cm[1][0]
FP=cm[0][1]
print('Testing Accuracy:',(TP+TN)/(TP+TN+FN+FP))
```

```
Testing Accuracy: 0.9230769230769231
```

```
[131]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,prediction1)
```

```
[131]: 0.9230769230769231
```

```
[132]: from sklearn.metrics import classification_report
print(classification_report(y_test, prediction1))
```

	precision	recall	f1-score	support
0	0.92	0.90	0.91	40
1	0.92	0.94	0.93	51
accuracy			0.92	91
macro avg	0.92	0.92	0.92	91
weighted avg	0.92	0.92	0.92	91

DECISION TREE

```
[134]: from sklearn.tree import DecisionTreeClassifier

dtc=DecisionTreeClassifier()
model2=dtc.fit(X_train,y_train)
prediction2=model2.predict(X_test)
cm2= confusion_matrix(y_test,prediction2)
```

```
[135]: cm2
```

```
[135]: array([[33,  7],
       [15, 36]], dtype=int64)
```

```
[136]: accuracy_score(y_test,prediction2)
```

```
[136]: 0.7582417582417582
```

```
[137]: print(classification_report(y_test, prediction2))
```

	precision	recall	f1-score	support
0	0.69	0.82	0.75	40
1	0.84	0.71	0.77	51
accuracy			0.76	91
macro avg	0.76	0.77	0.76	91
weighted avg	0.77	0.76	0.76	91

```
[139]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
```

```
[140]: from sklearn.svm import SVC
```

```

svm=SVC()
model4=svm.fit(X_train,y_train)
prediction4=model4.predict(X_test)
cm4= confusion_matrix(y_test,prediction4)

```

```
[141]: cm4
```

```
[141]: array([[33,  7],
           [ 2, 49]], dtype=int64)
```

```
[142]: accuracy_score(y_test, prediction4)
```

```
[142]: 0.9010989010989011
```

```
[143]: from sklearn.naive_bayes import GaussianNB
```

```

NB = GaussianNB()
model5 = NB.fit(X_train, y_train)
prediction5 = model5.predict(X_test)
cm5= confusion_matrix(y_test, prediction5)

```

```
[144]: cm5
```

```
[144]: array([[35,  5],
           [ 6, 45]], dtype=int64)
```

```
[145]: accuracy_score(y_test, prediction5)
```

```
[145]: 0.8791208791208791
```

```

[146]: print('cm4', cm4)
       print('-----')
       print('cm5',cm5)

```

```

cm4 [[33  7]
     [ 2 49]]
-----

```

```

cm5 [[35  5]
     [ 6 45]]

```

```
[147]: from sklearn.neighbors import KNeighborsClassifier
```

```

KNN = KNeighborsClassifier()
model6 = KNN.fit(X_train, y_train)
prediction6 = model6.predict(X_test)
cm6= confusion_matrix(y_test, prediction5)
cm6

```

```
[147]: array([[35,  5],  
              [ 6, 45]], dtype=int64)
```

```
[149]: print('KNN :', accuracy_score(y_test, prediction6))  
       print('LR :', accuracy_score(y_test, prediction1))  
       print('DT :', accuracy_score(y_test, prediction2))  
       print('NB: ', accuracy_score(y_test, prediction4))  
       print('SVM :', accuracy_score(y_test, prediction5))
```

KNN : 0.8351648351648352

LR : 0.9230769230769231

DT : 0.7582417582417582

NB: 0.9010989010989011

SVM : 0.8791208791208791

Best Accuracy: Logistic Regression : 92 Same Accuracy NB and Decision tree : 90